

Lab 2

DBsys

Elia Faure-Rolland
Tuomas Rahkola
Jonas Thalmeier

26/11/2024



Design Overview

In this lab, we implemented core modules for the SimpleDB system, focusing on managing heap files and sequential scanning. Key design choices included:

- **HeapFile and HeapPage Structure:** The heap file structure was designed to store pages sequentially on disk. Each page contains a header (bitmask) and fixed-size tuple slots. Header bits indicate whether corresponding tuple slots are occupied. Also in the HeapFile class we fixed the DEFAULT_PAGE_SIZE to 4096, in order to avoid calling BufferPool methods when reading from the disk.
- **Tuple and TupleDesc Management:** Tuples were defined with schemas provided by TupleDesc, which describes field names and types. These components enable consistent tuple manipulation across the database.
- **Buffer Pool for Page Management:** The BufferPool class caches pages, minimizing disk access. It retrieves pages using the getPage method, which checks for cached pages or reads from disk if unavailable.
- **Iterative Design in SeqScan:** The SeqScan operator was implemented to sequentially iterate over tuples in the database using the DbFileIterator.

Key Problems Encountered

Four major challenges arose during the implementation:

1. Hashing Function in HeapPageId

The initial implementation of the hashCode method in HeapPageId used a naive approach by concatenating table and page IDs as strings and parsing them into integers. This caused an **integer overflow**, leading to unexpected behavior when HeapPageId was used as a hash key. To resolve this, we adopted Java's built-in Objects.hash method:

```
return Objects.hash(tableId, pgNo);
```

This approach ensures a reliable hash code without risking overflow, leveraging robust hashing provided by the Java standard library.

2. Header Size Calculation in HeapPage

The getHeaderSize method initially used integer division:

```
return (int) Math.ceil(this.getNumTuples() / 8);
```

This caused an error in calculating the number of bytes required for the page header due to the truncation of integer division. The issue was corrected by casting the divisor to double, ensuring proper division:

```
return (int) Math.ceil(this.getNumTuples() / 8.0);
```

This adjustment ensures accurate header size computation, especially for pages with a small number of tuples.

3. Equals method in HeapPageId

The `Equals` method in `HeapPageId` originally used:

```
return this.tableId == other.tableId && this.pgNo == other.pgNo;
```

Because this did not correctly compare the two, it caused the `getPage()` method to not be able to find previous pages from cache. The issue was fixed by changing the `==` to `.equals()`, that compares the two correctly:

```
return this.tableId.equals(other.tableId) && this.pgNo.equals(other.pgNo);
```

The failure of `getPage` was noticed on the `ScanTest.testCache` test, which failed due to bufferpool, getting full.

4. Issue with `next()` Method in `HeapFile Iterator`

The original implementation of the `next()` method in the `HeapFile` iterator encountered an issue due to multiple calls to check for the next tuple within the method. This redundancy led to the page index being incremented twice under specific conditions, causing the iterator to skip pages unexpectedly. Surprisingly all the tests worked anyhow.

To prevent problems in the next labs, we simplified the method to ensure that the check for the next tuple occurs only once. This change prevents redundant increments and ensures proper tracking of the current page and tuple. As a result, the iterator now retrieves tuples sequentially and reliably without skipping pages.

Running the First Query

After resolving all encountered issues and bugs, the scan query code provided in the assignment was implemented as a class named `ScanQuery`, saved in a Java file of the same name within the `simplifiedb` directory. For testing purposes, a `.txt` file containing 100 rows and 4 columns of integers ranging from 1 to 10 was created. This text file was converted into a `.dat` file using the provided command. The query code was then executed, and the output successfully matched the input data.

Conclusions

The project emphasized the importance of careful consideration in method implementation, especially when relying on fundamental operations like hashing and division. Addressing these issues improved both correctness and robustness in handling database pages. Overall we spent about 20-30 hours combined in the span of 3 weeks. Most of it was spent implementing the code, but significant part was also spent on bug fixing. It was also hard to pinpoint the reason for the errors using the provided tests due to the recursive nature of the database.